

EXPERIMENTROADMAP

xOS

Validate Before You Build

Test the riskiest assumptions first. Reduce uncertainty with evidence, not opinions.

Based on David Bland's Testing Business Ideas framework (Desirability × Viability × Feasibility).

Bjarn Brunenberg

February 2026

CONFIDENTIAL

1. xOS Idea Scorecard

Before investing time and money, every business idea should be rigorously evaluated. Using a three-criteria framework weighted by importance, here is how xOS scores as a business opportunity.

1.1 Specificity (40% weight) — Score: 7.5/10

Problem clarity: Strong. CRO teams are expensive (€800K–€2M/year for 8–12 specialists), workflows are fragmented across 5+ tools, and knowledge doesn't compound. The problem is real and quantifiable.

Target audience: Good but needs sharpening. Two distinct segments: (1) mid-market to enterprise companies with 3+ CRO specialists, and (2) individual experimentation practitioners. The B2B segment is clearer and higher-value.

Solution approach: Clear and differentiated. An AI operating system built on Claude Code with specialized skills, MCP integrations, and a marketplace for expert brains. Nobody else is building this specific combination.

1.2 Focus (30% weight) — Score: 6.0/10

Core value proposition: "Replace 80% of manual CRO work" is compelling and measurable. Strong.

Scope risk: The vision spans 9 technical modules, a marketplace, a consulting arm, a dashboard, team and individual subscriptions, and GitHub distribution. This is a LOT for v1. The biggest risk is trying to do everything at once and delivering nothing well.

Recommendation: Narrow to ONE beachhead customer and ONE killer module for v1. The rest is roadmap, not launch.

1.3 Relevance (30% weight) — Score: 8.5/10

Market timing: Near-perfect. AI agents are mainstream conversation. Claude Code is gaining traction. Companies are asking "how do we use AI for growth?" xOS answers that question specifically for experimentation.

Problem significance: High. The promise of 80% cost reduction while maintaining velocity is exactly what CFOs want to hear.

Competitive landscape: No direct competitor. VWO and Optimizely are testing tools, not operating systems. CRO agencies sell labor, not intelligence. First-mover advantage is real.

Overall Score

Criterion	Weight	Score	Weighted
Specificity	40%	7.5	3.00
Focus	30%	6.0	1.80

Criterion	Weight	Score	Weighted
Relevance	30%	8.5	2.55
TOTAL	100%		7.35 / 10

VERDICT: 7.35/10

Strong idea with high market relevance and clear problem-solution fit. The main risk is scope: the vision is broad enough to become unfocused. Narrow the v1 scope aggressively and this scores 8+.

2. Business Model (Updated)

2.1 Revenue Streams

Stream	Price	Description
Base Subscription (xOS Brain)	€19/month	Core experimentation operating system. Full brain with hypothesis factory, experiment design, analysis engine, report generator. Hosted on private GitHub. Continuous updates, bug fixes, new MCPs.
Hero Brain Add-on	€29/month	Choose up to 5 “hero brains” from the marketplace. Access expert brains from Ton Wesseling (statistics), Erin Weigel (design), Lukas Vermeer (strategy), Kevin Anderson (career/benchmarking), and more.
Enterprise Implementation	€15–25K one-time	Bjarn implements xOS at the company. Sets up Claude Code, connects tools, trains the 2 human operators. Includes first 5 experiments guided.
Enterprise Team License	Custom pricing	Multi-seat license for teams. Volume discounts. Dedicated support. Custom integrations.

2.2 Revenue Distribution

Base subscription revenue goes entirely to xOS (Bjarn). Hero brain revenue is shared: the majority goes to xOS, with a percentage going to the contributing expert. Experts whose brains are used most get a higher percentage as an incentive — creating a virtuous cycle where the best content earns the most.

2.3 Two Sales Channels

Channel 1 — Self-Serve Subscriptions: Individual and team plans via private GitHub access. LinkedIn thought leadership as primary acquisition. Community-driven growth.

Channel 2 — Enterprise Implementation: High-touch consulting. Bjarn goes into companies, sets up xOS, trains the team. High-ticket (€15–25K) that also converts to ongoing subscriptions.

3. Riskiest Assumptions — D × V × F

Following David Bland's framework from Testing Business Ideas, every business idea rests on three types of assumptions. Desirability: do customers want this? Viability: can we make money? Feasibility: can we build and deliver it? The riskiest assumptions must be tested first.

3.1 Desirability — “Do they want this?”

ID	Assumption	Risk	Evidence Today	Why It Matters
D1	CRO teams actively suffer from fragmented tooling and want a unified operating system	HIGH	Anecdotal from Bjarn's own experience. No systematic customer evidence.	If teams are “fine” with their current setup, there's no compelling reason to buy.
D2	Companies will trust an AI system for strategic experimentation decisions	CRITICAL	Zero evidence. Single biggest unknown.	If companies don't trust AI output, xOS becomes a toy, not a tool.
D3	CRO managers will champion (not block) a tool that replaces 80% of their team's work	HIGH	No evidence. Turkeys don't vote for Christmas.	The buyer and user may have conflicting incentives. Head of CRO might see xOS as a threat.
D4	Individual practitioners will pay €19/mo for a pre-built brain vs. building their own	MEDIUM	Evo's Second Brain sold €50K+ in 2 months. Market proof exists.	If people prefer DIY, the subscription model fails for individuals.
D5	People want to “rent” a thought leader's brain vs. reading their content/courses	HIGH	No evidence. Completely new concept.	If hero brains don't feel better than free content + ChatGPT, nobody upgrades to €29.
D6	The addressable market is large enough to sustain the business	MEDIUM	Global A/B testing market ~\$1.6B. Thousands at experimentation conferences.	If the market is too niche, revenue caps out quickly.

3.2 Viability — “Can we make money?”

ID	Assumption	Risk	Evidence Today	Why It Matters
V1	€19/month generates enough subscribers for meaningful revenue	MEDIUM	SaaS benchmarks suggest this range works for individual tools.	Too low = leaves money on table. Too high = limits adoption.

ID	Assumption	Risk	Evidence Today	Why It Matters
V2	€29/month hero tier justifies the 50% price increase	MEDIUM	No evidence.	If hero brains aren't worth €10 more, upgrade rates near zero.
V3	Thought leaders will build and maintain their brains for a revenue share	HIGH	No commitments yet. Just Bjarn's idea.	If experts don't participate, the marketplace doesn't exist.
V4	Enterprise companies will pay €15–25K for implementation	MEDIUM	Consulting market validates willingness to pay for setup.	If companies learn Claude Code fast, implementation fee is hard to justify.
V5	Monthly churn stays below 5%	HIGH	No evidence. AI tool fatigue is real in 2025–2026.	If people try for a month and churn, the LTV doesn't support the business.

3.3 Feasibility — “Can we build it?”

ID	Assumption	Risk	Evidence Today	Why It Matters
F1	Claude Code's skill/MCP architecture reliably supports the full xOS workflow	MEDIUM	Bjarn's own brain works. Architecture proven at small scale.	If skills conflict or break at scale, the product fails.
F2	Non-technical CRO professionals can install, configure, and use Claude Code daily	CRITICAL	Bjarn's community has “very low” Claude Code literacy.	If the target user can't use the tool, nothing else matters. Potential dealbreaker.
F3	The “secret sauce” can be protected when packaged via MCP	HIGH	Mateusz working on MCP packaging. Not yet proven.	If buyers can see/copy the skills, there's no moat and no recurring revenue.
F4	AI brain produces consistently reliable analysis (not hallucinated stats)	HIGH	LLMs struggle with statistical reasoning. No evidence xOS is different.	If analysis is wrong even 10% of the time, trust collapses.
F5	GitHub-based distribution with subscription access control works at scale	MEDIUM	GitHub private repos with team access is established.	Operational headache if subscription management becomes manual.

4. Prioritized Test Plan

Not all assumptions are equal. Some are critical (if wrong, the entire idea fails). The table below ranks all assumptions by risk and maps each to a specific experiment from David Bland's Testing Business Ideas.

#	ID	Assumption	Bucket	Risk	Test Method (Bland)	Timeline
1	D2	Companies trust AI for experimentation decisions	Desir.	CRITICAL	Customer Interviews (p.106) + Expert Interviews (p.115)	Week 1
2	F2	Non-technical CROs can use Claude Code	Feas.	CRITICAL	Concierge (p.248) – walk 3 users through xOS manually	Week 2–3
3	D3	CRO managers won't block as job threat	Desir.	HIGH	Customer Interviews (p.106) with CRO managers specifically	Week 1
4	D1	Real pain from fragmented CRO tooling	Desir.	HIGH	Discovery Survey (p.122) + Online Ad (p.146)	Week 1–2
5	D5	People want to rent thought leader brains	Desir.	HIGH	Simple Landing Page (p.260) – measure sign-up intent	Week 2
6	V3	Thought leaders will participate for rev share	Viab.	HIGH	Partner Interviews (p.114) with Ton, Erin, Lukas, Kevin	Week 2
7	F4	AI produces reliable statistical analysis	Feas.	HIGH	Data Sheet (p.190) – test on 10 past experiments	Week 3–4
8	V5	Monthly churn stays below 5%	Viab.	HIGH	Track from day 1. Set up cohort analysis.	Post-launch
9	F3	IP protection via MCP works	Feas.	HIGH	Partner Interview with Mateusz (p.114)	Week 2
10	D4	Practitioners pay €19/mo vs. DIY	Desir.	MEDIUM	Presale (p.274) – early access at €19/mo	Week 3
11	V1	€19/mo is the right price point	Viab.	MEDIUM	Online Ad (p.146) – test different price points	Week 2–3
12	D6	Market is large enough	Desir.	MEDIUM	Search Trend Analysis (p.126) + Web Traffic Analysis (p.130)	Week 1

#	ID	Assumption	Bucket	Risk	Test Method (Bland)	Timeline
13	F1	Claude Code architecture supports full workflow	Feas.	MEDIUM	Feature Stub (p.156) — build one end-to-end flow	Week 3–4
14	V4	Companies pay €15–25K for implementation	Viab.	MEDIUM	Letter of Intent (p.294) — 2–3 company commitments	Week 4

5. The 4-Week Pre-MVP Validation Sprint

Before going to Brazil and investing development time, run this 4-week sprint to gather evidence on your riskiest assumptions. Total budget: under €500. If all goes well, you'll have enough conviction (and early revenue signals) to greenlight development.

Week 1: Discovery — Talk to People

Goal: Understand if the pain is real and if people trust AI for experimentation.

Day	Activity	Method (Bland)	Target	Success Metric
Mon–Tue	5 Customer Interviews with CRO professionals	Customer Interviews (p.106)	CRO specialists, managers	4/5 confirm fragmentation pain AND openness to AI
Wed	3 Expert Interviews with CRO leaders	Expert Interviews (p.115)	Senior CROs	2/3 would consider an AI operating system
Thu	Market sizing analysis	Search Trend + Web Traffic Analysis	N/A	Addressable market > 10,000 potential users
Fri	Synthesize learnings. Go/No-Go on Week 2.	Team decision	N/A	Clear signal on D1, D2, D3

Week 2: Signal Testing — Attract Interest

Goal: Test whether the value proposition resonates and whether partners are interested.

Day	Activity	Method (Bland)	Target	Success Metric
Mon	Build xOS landing page (Carrd or similar) with email capture	Simple Landing Page (p.260)	Public	Live page with clear value proposition
Tue–Wed	Run LinkedIn + targeted ads (€100–200 budget)	Online Ad (p.146)	CRO professionals	> 5% CTR, > 15% email conversion
Thu	Call Ton Wesseling, Erin, Lukas, Kevin about hero brains	Partner Interviews (p.114)	Thought leaders	2+ express genuine interest in contributing
Thu	Call Mateusz about MCP packaging / IP protection	Partner Interview (p.114)	Second Brain Factory	Clear understanding of packaging feasibility
Fri	Synthesize. Count sign-ups. Go/No-Go on Week 3.	Team decision	N/A	50+ email sign-ups, 2+ partner commits

Week 3: Validation — Can People Use It?

Goal: Test usability with non-technical CROs and validate willingness to pay.

Day	Activity	Method (Bland)	Target	Success Metric
Mon–Tue	Build minimal concierge: one skill (hypothesis factory)	Feature Stub (p.156)	N/A	Working hypothesis-to-output flow
Wed–Thu	Walk 3 CRO professionals through using it (screen share)	Concierge (p.248)	3 non-technical CROs	2/3 complete flow with minimal hand-holding
Thu	Send presale offer to email list at €19/mo early bird	Presale (p.274)	Email list	5+ presales (10% of list)
Fri	Share €29/mo hero tier concept, measure upgrade interest	Landing Page variant	Email list	30%+ of €19 buyers interested in €29

Week 4: Decision Sprint — Build or Pivot?

Goal: Compile all evidence and make the final go/no-go decision.

Day	Activity	Details
Mon–Tue	Test AI quality on 5 real past experiments	Compare xOS analysis output to actual human analyst results. Score accuracy.
Wed	Compile evidence dashboard	Single-page evidence summary from all metrics across Weeks 1–3.
Thu	Enterprise outreach: 2–3 companies for Letter of Intent	Gauge willingness to pay €15–25K for implementation.
Fri	FINAL DECISION: Build / Pivot / Kill	Use Go/No-Go criteria (Section 6). Be honest about the evidence.

6. Go / No-Go Decision Criteria

Use this table at the end of Week 4 to make a rigorous decision. Green = proceed to build. Yellow = adjust and retest. Red = kill or major pivot.

Metric	Green (Go)	Yellow (Pivot)	Red (Kill)
Customer interviews confirm pain (D1)	4/5+ confirm	3/5 confirm	<3/5 confirm
Trust in AI for experimentation (D2)	Openness + curiosity	Skeptical but willing	Strong resistance
CRO manager reaction (D3)	See it as empowering	Cautious, not hostile	Active resistance
Landing page email conversions	> 15% conversion rate	8–15% conversion	< 8% conversion
Total email sign-ups (€200 spend)	50+ sign-ups	25–50 sign-ups	< 25 sign-ups
Presale conversions at €19/mo	5+ paying (10%+)	2–4 paying	<2 paying
Hero brain partner interest	3+ thought leaders	1–2 interested	0 committed
Concierge usability test	2/3 can use it	1/3 can use it	0/3 can use it
AI analysis quality vs human	85%+ accuracy	70–85% accuracy	< 70% accuracy
Enterprise interest (LOI)	2+ companies	1 company	0 interested

REMEMBER

You are an experimenter at heart. Eat your own dog food. Run this validation sprint the same way you'd run an experimentation program for a client. Be rigorous. Be honest about the evidence. The goal is not to confirm your excitement – it's to reduce uncertainty so you can invest with confidence.

7. What to Bring to Your Developer in Brazil

If the evidence from Weeks 1–4 gives you a Green or strong Yellow signal, here is what you bring to your developer:

- **The Product Vision document (already created)**
- **This Experiment Roadmap with all evidence collected**
- **A clear decision on v1 scope:** which 2–3 modules to build first (likely: Hypothesis Factory + Analysis Engine + Report Generator)
- **Your CLAUDE.md master brain draft** — your experimentation philosophy encoded
- **A list of MCPs needed:** Airtable, VWO, Google Sheets
- **Evidence dashboard** showing customer demand signals
- **5+ presale customers** as proof of demand
- **Partner commitments** from thought leaders for the marketplace

Test first. Build with evidence. Scale with conviction.

This is what xOS teaches companies to do — now do it yourself.